

Reinforcement learning

Episode 5

Deep reinforcement learning

What we already know:

- Q-learning

$$L = E_{s \sim S, a \sim A} [(Q(s, a) - (r + \gamma \max_{a'} Q(s', a'))))^2]$$

- Approximation of q-values with respect to state $Q(s, a, \Theta)$, where Θ is the vector of weights
- Experience replay

This is not enough!

Autocorrelation

- Target is based on prediction

$$r + \gamma \max_{a'} Q(s', a', \Theta)$$

- Since we use function approximation, when we update $Q(s, a, \Theta)$ we also update $Q(s', a, \Theta)$ towards that direction
- In worst case network may diverge, but usually it becomes unstable.
- **How to stabilize weights?**

Target network

Idea: use network with frozen weights to compute the target

$$L(\Theta) = E_{s \sim S, a \sim A} [(Q(s, a, \Theta) - (r + \gamma \max_{a'} Q(s', a', \Theta^-)))^2]$$

where Θ^- is the frozen weights

Hard target network:

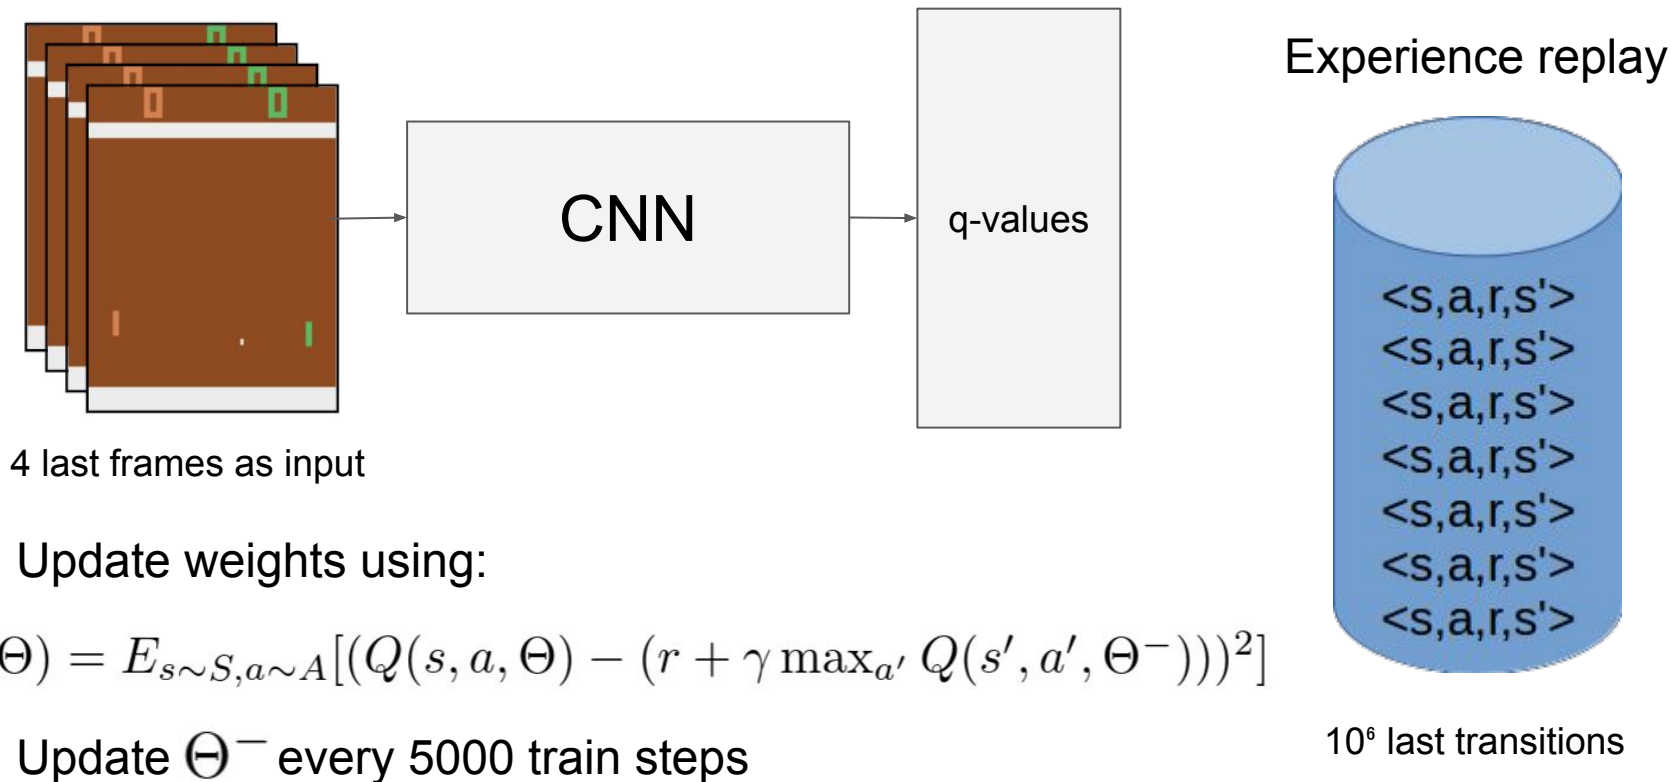
Update Θ^- every **n** steps and set its weights as Θ

Soft target network:

Update Θ^- every step:

$$\Theta^- = (1 - \alpha)\Theta^- + \alpha\Theta$$

Playing Atari with Deep Reinforcement Learning (2013, Deepmind)



Problem of overestimation

We use “max” operator to compute the target

$$L(s, a) = (Q(s, a) - (r + \gamma \max_{a'} Q(s', a')))^2$$


Surprisingly here is a problem

(although we want $E_{s \sim S, a \sim A}[L(s, a)]$ to be equal zero)

Problem of overestimation

Normal distribution
 $3 \cdot 10^6$ samples

mean: ~ 0.0004



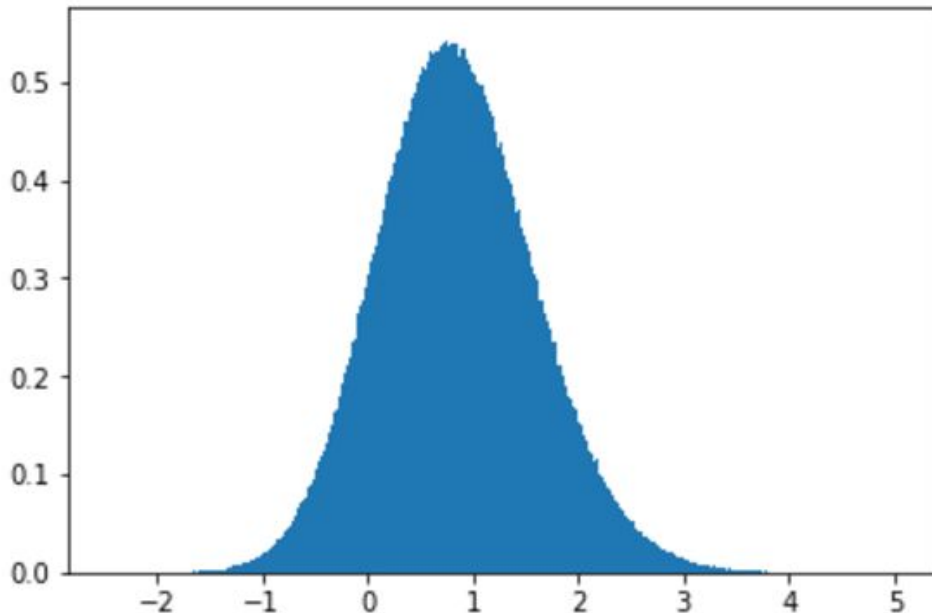
Problem of overestimation

Normal distribution

$3 \cdot 10^6 \times 3$ samples

Then take maximum of every tuple

mean: ~ 0.8467




Problem of overestimation

Normal distribution


$3 \times 10^6 \times 10$ samples

Then take maximum of every tuple

mean: ~ 1.538



Problem of overestimation



Suppose true $Q(s', a')$ are equal 0 for all a'

But we have an approximation (or other kind) error $\sim N(0, \sigma^2)$

So $Q(s, a)$ should be equal r

But if we update $Q(s, a)$ towards $r + \gamma \max_{a'} Q(s', a')$

we will have overestimated $Q(s, a) > r$ because

$$E[\max_{a'} Q(s', a')] \geq \max_{a'} E[Q(s', a')]$$

Double Q-learning (NIPS 2010)

Idea: use two estimators of q-values: Q^A, Q^B

They should compensate mistakes of each other because they will be independent

Let's get argmax from another estimator!

$$y = r + \gamma \max_{a'} Q(s', a') \quad \text{- Q-learning target}$$

$$y = r + \gamma Q(s', \operatorname{argmax}_{a'} Q(s', a')) \quad \text{- Rewritten Q-learning target}$$

$$y = r + \gamma Q^A(s', \operatorname{argmax}_a Q^B(s', a')) \quad \text{- Double Q-learning target}$$

Double Q-learning (NIPS 2010)

Algorithm 1 Double Q-learning

```
1: Initialize  $Q^A, Q^B, s$ 
2: repeat
3:   Choose  $a$ , based on  $Q^A(s, \cdot)$  and  $Q^B(s, \cdot)$ , observe  $r, s'$ 
4:   Choose (e.g. random) either UPDATE(A) or UPDATE(B)
5:   if UPDATE(A) then
6:     Define  $a^* = \arg \max_a Q^A(s', a)$ 
7:      $Q^A(s, a) \leftarrow Q^A(s, a) + \alpha(s, a) (r + \gamma Q^B(s', a^*) - Q^A(s, a))$ 
8:   else if UPDATE(B) then
9:     Define  $b^* = \arg \max_a Q^B(s', a)$ 
10:     $Q^B(s, a) \leftarrow Q^B(s, a) + \alpha(s, a) (r + \gamma Q^A(s', b^*) - Q^B(s, a))$ 
11:  end if
12:   $s \leftarrow s'$ 
13: until end
```


How to apply this algorithm in deep reinforcement learning?

Deep Reinforcement Learning with Double Q-learning (Deepmind, 2015)

Idea: use main network to choose action!

$$y_{dq_n} = r + \gamma \max_{a'} Q(s', a', \Theta^-)$$

$$y_{ddq_n} = r + \gamma Q(s', \operatorname{argmax}_{a'} Q(s', a', \Theta), \Theta^-)$$



	DQN	Double DQN	Double DQN (tuned)
Median	47.5%	88.4%	116.7%
Mean	122.0%	273.1%	475.2%

Prioritized Experience Replay (2016, Deepmind)

Idea: sample transitions from xp-replay more clever

We want to set probability for every transition. Let's use the absolute value of TD-error of transition as a probability!

$$\text{TD-error } \delta = Q(s, a) - (r + \gamma Q(s', \operatorname{argmax}_{a'} Q(s', a', \Theta), \Theta^-))$$

$$p = |\delta|$$

$$P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha} \text{ where } \alpha \text{ is the priority parameter (when } \alpha \text{ is 0 it's the uniform case)}$$

Do you see the problem?

Prioritized Experience Replay (2016, Deepmind)

Idea: sample transitions from xp-replay more clever

We want to set probability for every transition. Let's use the absolute value of TD-error of transition as a probability!

$$\text{TD-error } \delta = Q(s, a) - (r + \gamma Q(s', \arg\max_{a'} Q(s', a', \Theta), \Theta^-))$$

$$p = |\delta|$$

$$P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha} \text{ where } \alpha \text{ is the priority parameter (when } \alpha \text{ is 0 it's the uniform case)}$$

Do you see the problem?

Transitions become non i.i.d. and therefore we introduce the bias,

Prioritized Experience Replay (2016, Deepmind)

Solution: we can correct the bias by using importance-sampling weights

$$w_i = \left(\frac{1}{N} \cdot \frac{1}{P(i)} \right)^\beta \quad \text{where } \beta \text{ is the parameter}$$

We also normalize weights by $1 / \max_i w_i$ (here is not mathematical reason)

When we put transition into experience replay, we set maximal priority $p_t = \max_{i < t} p_i$

Prioritized Experience Replay (2016, Deepmind)

Algorithm 1 Double DQN with proportional prioritization

```
1: Input: minibatch  $k$ , step-size  $\eta$ , replay period  $K$  and size  $N$ , exponents  $\alpha$  and  $\beta$ , budget  $T$ .
2: Initialize replay memory  $\mathcal{H} = \emptyset$ ,  $\Delta = 0$ ,  $p_1 = 1$ 
3: Observe  $S_0$  and choose  $A_0 \sim \pi_\theta(S_0)$ 
4: for  $t = 1$  to  $T$  do
5:   Observe  $S_t, R_t, \gamma_t$ 
6:   Store transition  $(S_{t-1}, A_{t-1}, R_t, \gamma_t, S_t)$  in  $\mathcal{H}$  with maximal priority  $p_t = \max_{i < t} p_i$ 
7:   if  $t \equiv 0 \pmod K$  then
8:     for  $j = 1$  to  $k$  do
9:       Sample transition  $j \sim P(j) = p_j^\alpha / \sum_i p_i^\alpha$ 
10:      Compute importance-sampling weight  $w_j = (N \cdot P(j))^{-\beta} / \max_i w_i$ 
11:      Compute TD-error  $\delta_j = R_j + \gamma_j Q_{\text{target}}(S_j, \arg \max_a Q(S_j, a)) - Q(S_{j-1}, A_{j-1})$ 
12:      Update transition priority  $p_j \leftarrow |\delta_j|$ 
13:      Accumulate weight-change  $\Delta \leftarrow \Delta + w_j \cdot \delta_j \cdot \nabla_\theta Q(S_{j-1}, A_{j-1})$ 
14:    end for
15:    Update weights  $\theta \leftarrow \theta + \eta \cdot \Delta$ , reset  $\Delta = 0$ 
16:    From time to time copy weights into target network  $\theta_{\text{target}} \leftarrow \theta$ 
17:  end if
18:  Choose action  $A_t \sim \pi_\theta(S_t)$ 
19: end for
```


Prioritized Experience Replay (2016, Deepmind)

Algorithm 1 Double DQN with proportional prioritization

```
1: Input: minibatch  $k$ , step-size  $\eta$ , replay period  $K$  and size  $N$ , exponents  $\alpha$  and  $\beta$ , budget  $T$ .
2: Initialize replay memory  $\mathcal{H} = \emptyset$ ,  $\Delta = 0$ ,  $p_1 = 1$ 
3: Observe  $S_0$  and choose  $A_0 \sim \pi_\theta(S_0)$ 
4: for  $t = 1$  to  $T$  do
5:   Observe  $S_t, R_t, \gamma_t$ 
6:   Store transition  $(S_{t-1}, A_{t-1}, R_t, \gamma_t, S_t)$  in  $\mathcal{H}$  with maximal priority  $p_t = \max_{i < t} p_i$ 
7:   if  $t \equiv 0 \pmod K$  then
8:     for  $j = 1$  to  $k$  do
9:       Sample transition  $j \sim P(j) = p_j^\alpha / \sum_i p_i^\alpha$ 
10:      Compute importance-sampling weight  $w_j = (N \cdot P(j))^{-\beta} / \max_i w_i$ 
11:      Compute TD-error  $\delta_j = R_j + \gamma_j Q_{\text{target}}(S_j, \arg \max_a Q(S_j, a)) - Q(S_{j-1}, A_{j-1})$ 
12:      Update transition priority  $p_j \leftarrow |\delta_j|$ 
13:      Accumulate weight-change  $\Delta \leftarrow \Delta + w_j \cdot \delta_j \cdot \nabla_\theta Q(S_{j-1}, A_{j-1})$ 
14:    end for
15:    Update weights  $\theta \leftarrow \theta + \eta \cdot \Delta$ , reset  $\Delta = 0$ 
16:    From time to time copy weights into target network  $\theta_{\text{target}} \leftarrow \theta$ 
17:  end if
18:  Choose action  $A_t \sim \pi_\theta(S_t)$ 
19: end for
```

It is the bonus homework!

Bootstrapped DQN (2016, Deepmind & Stanford)



Idea: make exploration ~~great again~~ more clever.

Usually we use e-greedy strategy to pick action.

Now we maintain K separate heads (for example, FC-layers) and shared body (for example, convolutional) weights.

Every episode we pick a head randomly and train both this head and shared network on transitions from that episode.

It allow us to make **deep** exploration.

Often to explore agent should make several non-greedy actions in a row, e-greedy strategy doesn't allow to do it.

Let's watch a video...

<https://www.youtube.com/watch?v=UXurvvdY93o>


Dueling Network Architectures for Deep Reinforcement Learning (2016, Deepmind)

Idea: change the network's architecture.

Recall:

Advantage Function $A(s,a) = Q(s,a) - V(s)$

So, $Q(s,a) = A(s,a) + V(s)$





Do you see the problem?

Dueling Network Architectures for Deep Reinforcement Learning (2016, Deepmind)

Here is one extra freedom degree!

Example:



Dueling Network Architectures for Deep Reinforcement Learning (2016, Deepmind)

Solution: require $\max_{a' \in |\mathcal{A}|} A(s, a'; \theta, \alpha)$ to be equal to zero!

So the **Q-function** computes as:

$$Q(s, a; \theta, \alpha, \beta) = V(s; \theta, \beta) + \left(A(s, a; \theta, \alpha) - \max_{a' \in |\mathcal{A}|} A(s, a'; \theta, \alpha) \right)$$

Authors of this papers also introduced this way to compute Q-values:

$$Q(s, a; \theta, \alpha, \beta) = V(s; \theta, \beta) + \left(A(s, a; \theta, \alpha) - \frac{1}{|\mathcal{A}|} \sum_{a'} A(s, a'; \theta, \alpha) \right)$$

They wrote that this variant increases stability of the optimization
(The fact that this loses the original semantics of Q doesn't matter)

Dueling Network Architectures for Deep Reinforcement Learning (2016, Deepmind)

Solution: require $\max_{a' \in |\mathcal{A}|} A(s, a'; \theta, \alpha)$ to be equal to zero!

So the **Q-function** computes as:

$$Q(s, a; \theta, \alpha, \beta) = V(s; \theta, \beta) +$$

$$\left(A(s, a; \theta, \alpha) - \max_{a' \in |\mathcal{A}|} A(s, a'; \theta, \alpha) \right)$$

It's the homework!

Authors of this papers also introduced this way to compute Q-values:

$$Q(s, a; \theta, \alpha, \beta) = V(s; \theta, \beta) +$$

$$\left(A(s, a; \theta, \alpha) - \frac{1}{|\mathcal{A}|} \sum_{a'} A(s, a'; \theta, \alpha) \right)$$

They wrote that this variant increases stability of the optimization
(The fact that this loses the original semantics of Q doesn't matter)

**LEARNING TO PLAY IN A DAY:
FASTER DEEP REINFORCEMENT LEARNING BY OPTIMALITY TIGHTENING(2016)**

Bellman equation:

$$Q(s, a) = r + \gamma \max_{a'} Q(s', a')$$

We can estimate it this way:

$$Q(s_t, a_t) = E[r_t + \gamma \max_{a_{t+1}} Q(s_{t+1}, a_{t+1})] \geq E[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots + \gamma^k \max_{a_{t+k}} Q(s_{t+k}, a_{t+k})]$$

Also we can rewrite it this way (divide all parts by γ^k and subtract k from indices)

$$Q(s_t, a_t) \leq \max_{a_t} Q(s_t, a_t) \leq E[\gamma^{-k} Q(s_{t-k}, a_{t-k}) - \gamma^{-k} r_{t-k} - \gamma^{-(k-1)} r_{t-k+1} - \dots - \gamma^{-1} r_{t-1}]$$

We can store these two estimations together with the transition:

$$U^{min} = \gamma^{-k} Q(s_{t-k}, a_{t-k}) - \gamma^{-k} r_{t-k} - \gamma^{-(k-1)} r_{t-k+1} - \dots - \gamma^{-1} r_{t-1}$$

$$L^{max} = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots + \gamma^k \max_{a_{t+k}} Q(s_{t+k}, a_{t+k})$$

**LEARNING TO PLAY IN A DAY:
FASTER DEEP REINFORCEMENT LEARNING BY OPTIMALITY TIGHTENING(2016)**


And now we can modify our loss function using these estimations:

$$y = r + \gamma Q(s', \operatorname{argmax}_{a'} Q(s', a', \Theta), \Theta^-)$$

$$L(\Theta) = E[(Q(s, a, \Theta) - y)^2 + \lambda(L^{max} - Q(s, a, \Theta))_+^2 + \lambda(Q(s, a, \Theta) - U_{min})_+^2]$$

Ours (10M)	less than 1 day (1 GPU)	345.70%	105.74%
DQN (200M)	more than 10 days (1 GPU)	241.06%	93.52%

Asynchronous Methods for Deep Reinforcement Learning (2016, Deepmind)



Questions?